

MyMONIT

Collecting measurements to monitor CERN's
experiments

Table of Contents

Table of Contents

| | |
|----------------------------------|----|
| High-level description..... | 2 |
| Requirements..... | 3 |
| Functional requirements..... | 4 |
| Non-Functional Requirements..... | 4 |
| Assumptions..... | 5 |
| Architecture..... | 6 |
| MyMONIT..... | 7 |
| Information flow..... | 8 |
| Security..... | 10 |
| Overview..... | 10 |
| Authentication..... | 10 |
| Authorization..... | 10 |
| Security Risks..... | 10 |
| Spoofing..... | 10 |
| Tampering..... | 11 |
| Repudiation..... | 11 |
| Information disclosure..... | 11 |
| Denial of service..... | 12 |
| Elevation of privilege..... | 12 |
| System Requirements..... | 14 |
| Storage space..... | 14 |
| CPU and memory..... | 14 |
| GDPR Consideration..... | 15 |
| References..... | 16 |
| Bibliography (TBD remove)..... | 17 |

Word count: 915

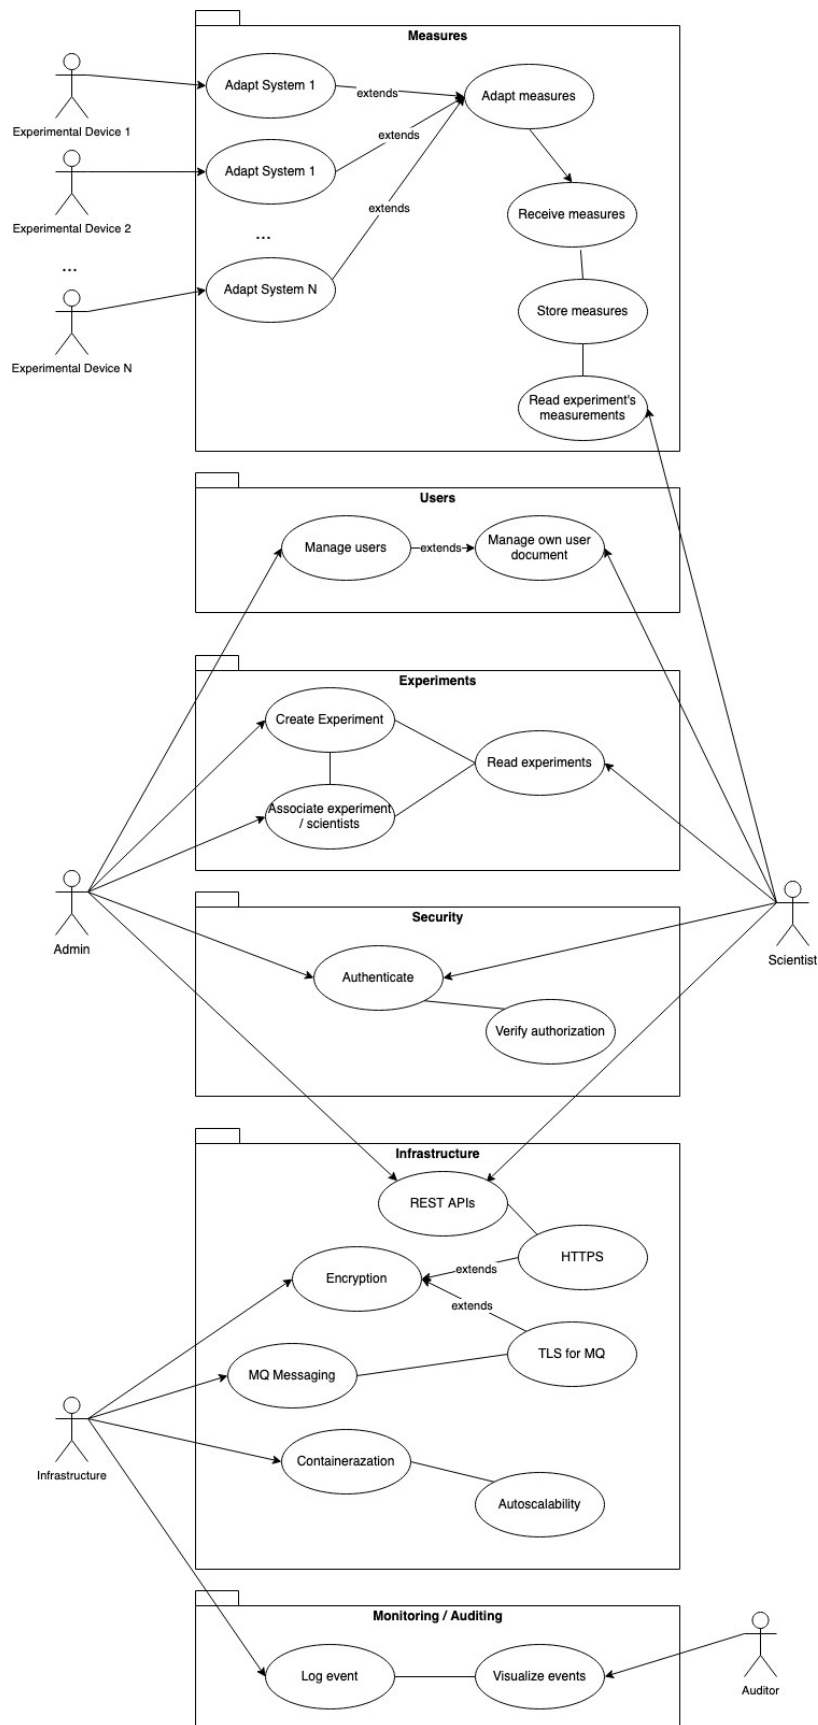
High-level description

CERN uses a variety of independently developed systems to monitor its infrastructure (Aimar et al., 2019). MyMONIT will be a solution to unify the monitoring of experiments into a single software integrating different streams of measurements to centralize this information.

MyMONIT will be scalable to ensure that it can cope with increasing demand. The solution will also include monitoring to detect anomalies in the system itself and the flow of the measurements.

Requirements

The following diagram illustrates all the use cases.



Functional requirements

- There will be three user types with the following role matrix:

| role | resource | scope | access |
|----------------|--------------|---|--------|
| Administrators | users | complete | RW |
| | experiments | complete | RW |
| | measurements | complete | R |
| | audits | No access | / |
| Scientists | users | user's record | RW |
| | experiments | only records associated with user | RW |
| | measurements | only records associated with user's experiments | R |
| | audits | No access | / |
| Auditors | users | No access | / |
| | experiments | No access | / |
| | measurements | No access | / |
| | audits | complete | R |

- For each source, an adapter will normalize the measure and transmit it to MyMONIT.
- The measures will be persisted, indexed per experiment, and made available through APIs to authorized scientists.
- A complete audit will be available from a separate interface.

Non-Functional Requirements

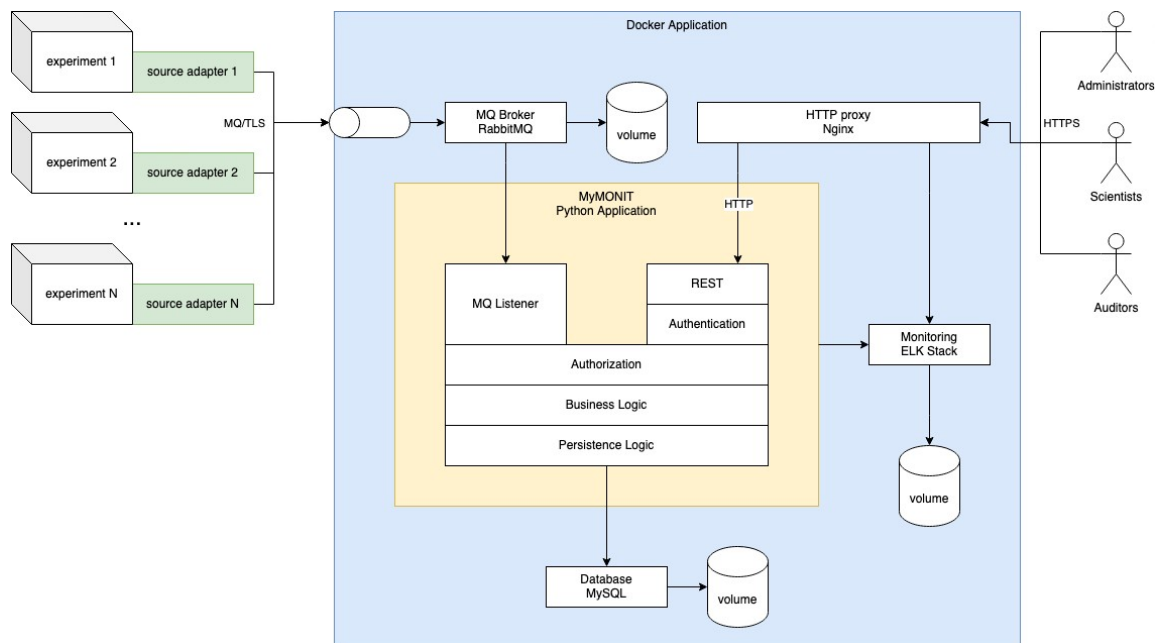
- Access points will be authenticated.
- Access to experiments will require per-user authorization.
- The system must serve concurrent users and concurrent experiments.

- 100% of data must be retained.

Assumptions

- The system's capacity must accommodate at least 10 years of data.
- Except for the database, all components of the solution will be stateless and allow for horizontal scalability. Autoscale functionalities in the clustering infrastructure will be sufficient to regulate the number of running instances and deal with variable demand (Kubernetes, N.A).
- It is expected an elevated flow of data coming from queues while users will visualize them almost in real-time. It is expected that the highest pressure could come from the input flow. Since reading a large amount of data could put pressure on the database, queries should be paginated.

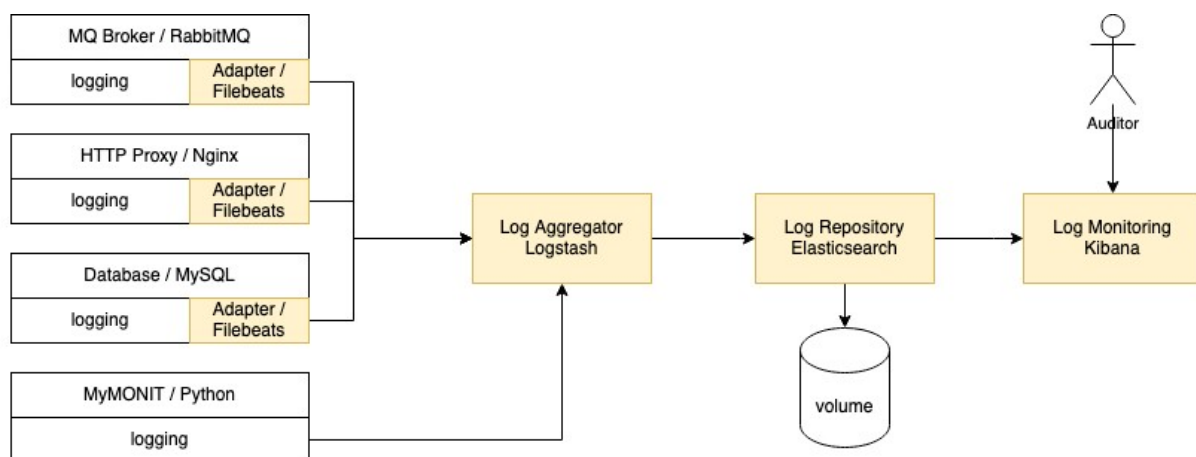
Architecture



The the adapters (in green) will send the measurements to the solution (in blue) where the main component (in yellow) will index them and expose them via REST APIs.

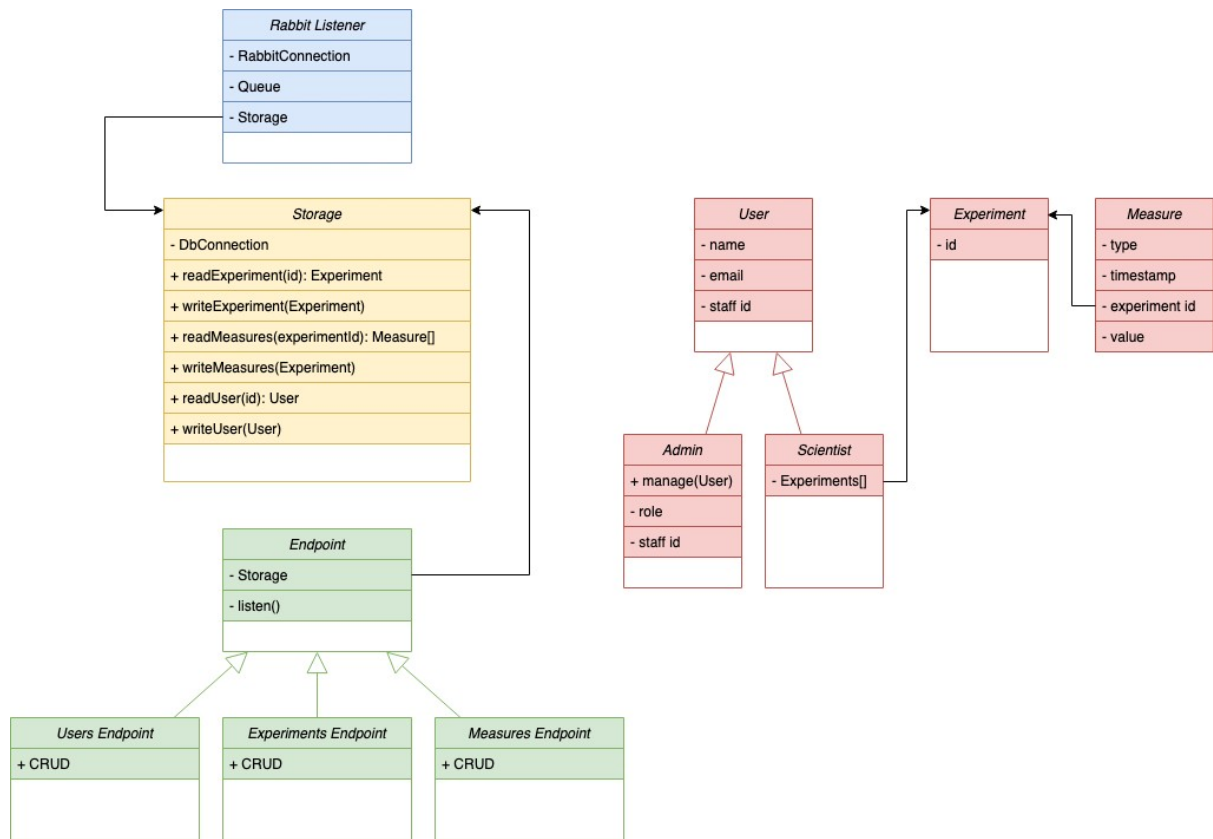
- Docker and Docker Compose: the solution will be containerized and will be deployable in compatible solutions such as Kubernetes.
- The adapters will be Python scripts customized to each specific case installed at the experiment's location.
- Nginx will be used as reverse proxy with SSL offloading and will hide all HTTP resources from the outside network.
- RabbitMQ, one of the most popular solutions in the market (Souza, 2020) will be used as MQ Broker to accept data streams from the experiments encrypted in TLS.

- MySQL will be responsible for the storage of the application's data. A SQL database was preferred for its simplicity. The design will allow to replace it with a NoSQL database in case of scalability issues (Khasawneh, 2020).
- MyMONIT will be a Python application using Flask and Pika.
- ELK Stack (Elastic Search, Logstash, and Kibana) will collect all logs and expose a dashboard for auditing. Filebeats will be used as an adapter where needed.



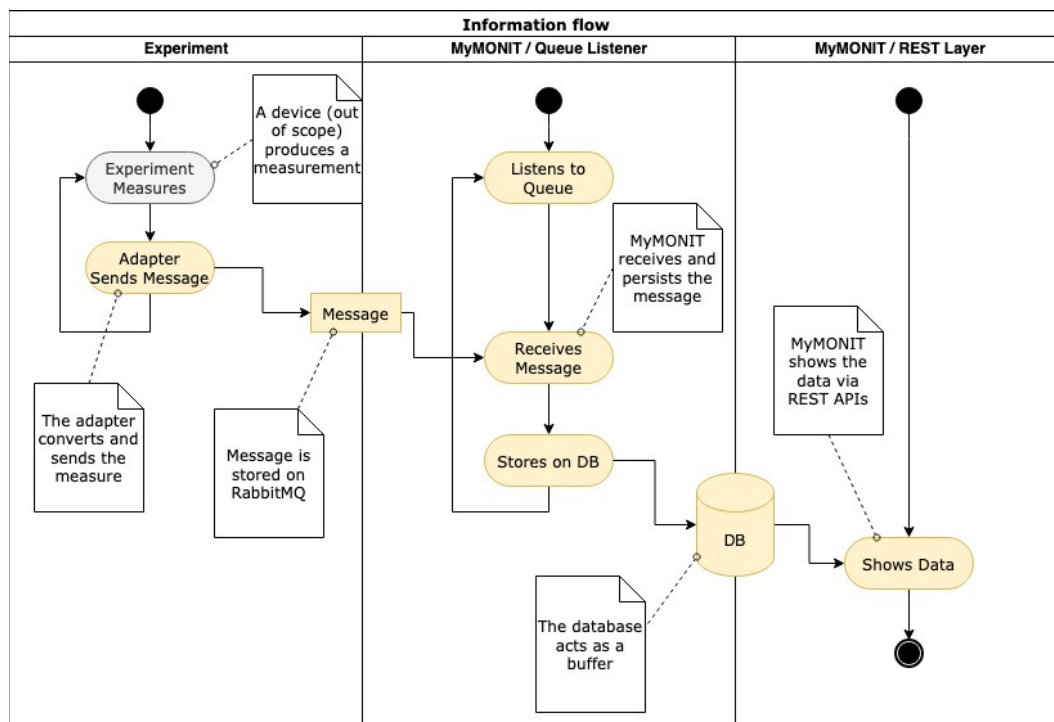
MyMONIT

There will be no direct interactions between the components consuming messages from the broker (in blue) and the components exposing REST endpoints (in green). The Storage (in yellow) will mediate the communications between the two parts.

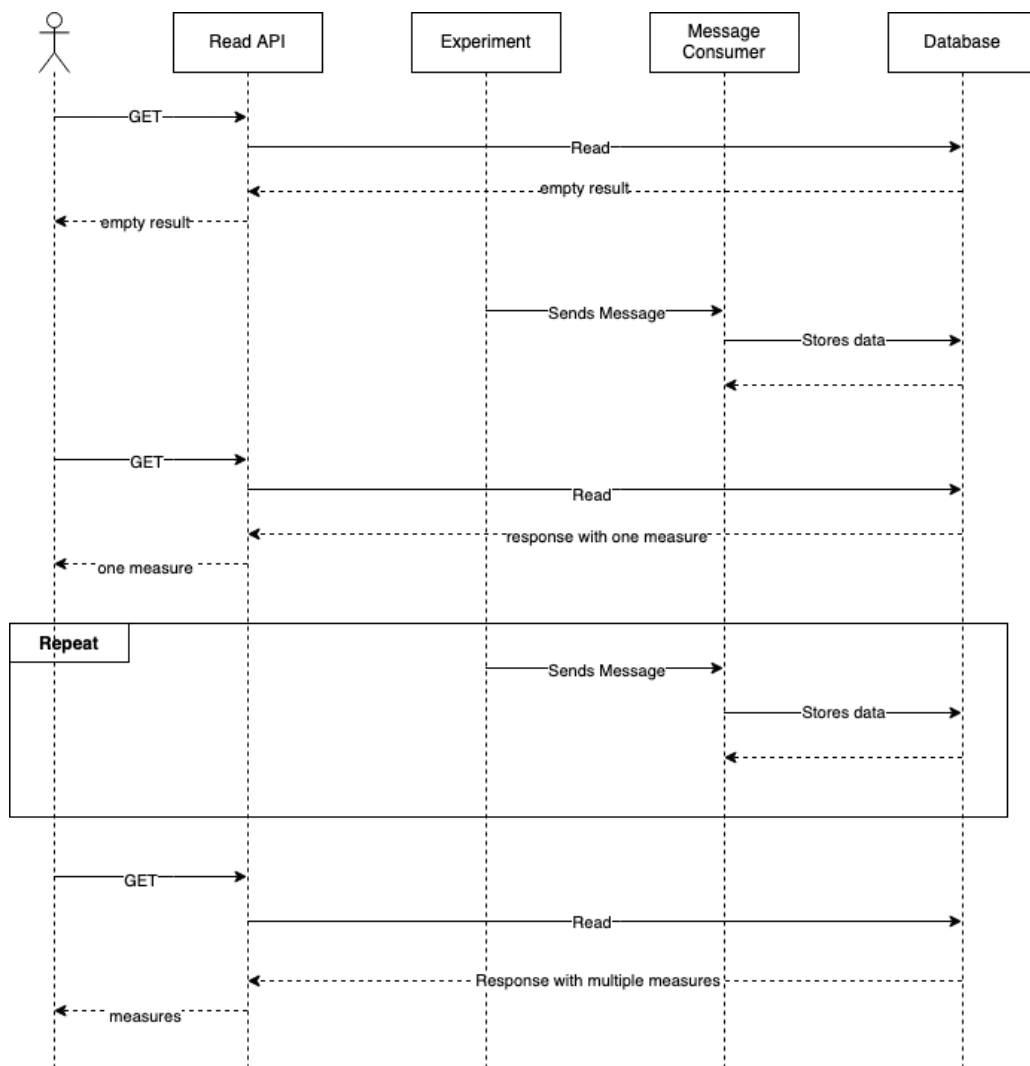


Information flow

The following diagram shows from a location perspective how information flows between components.



The following represents the same flow from a timeline perspective



Security

Overview

The main security concerns are the risks of sabotage and information leak. Being a monitoring tool, an attacker may try to disrupt its operations to cover the main attack on an experimental facility. Information leaks could endanger the process of peer reviews allowing scientists to steal data from parallel research.

Authentication

The authentication endpoint will validate credentials by comparing the input with the hashes stored in the database. The endpoint will return a JSON web token that will remain valid for a limited time. The token will be required in the calls to all other APIs. The authentication endpoint will require the inclusion of a shared secret (API key) in the request. This additional measure will limit the chances to perform a brute force attack (OWASP, N.D.).

Authorization

Authorization to the users, experiments, and measurements endpoints will be role-based. Auditors will have full access limited to audits.

Security Risks

Using the STRIDE model, the following threats were identified and classified with DREAD.

Spoofing

User's credentials violation

| Type | Level |
|------|-------|
|------|-------|

| | |
|-----------------|---|
| Damage | High, experiments would be exposed, users' records compromised, data leak |
| Reproducibility | High |
| Exploitability | High |
| Affected users | One user. All, if the user is administrator |
| Discoverability | Medium. User's credentials may be easy to guess |

Tampering

Introducing fake measurements on the message broker

| Type | Level |
|-----------------|---|
| Damage | High, experiments would be invalidated |
| Reproducibility | Medium. The highest risk is broker's authentication |
| Exploitability | High. Discovering credentials would make it easy to exploit the vulnerability |
| Affected users | All scientists |
| Discoverability | Medium. The broker is public, but credentials are highly secure |

Repudiation

Administrator denies a sabotage

| Type | Level |
|-----------------|--|
| Damage | High, experiments would be invalidated |
| Reproducibility | Low. All actions are audited. Administrator do not have W access to audits |
| Exploitability | Low. Administrators do not have W access to audits |
| Affected users | All scientists |
| Discoverability | Low. It requires the discovery of additional vulnerabilities |

Information disclosure

Database breach

| Type | Level |
|-----------------|---|
| Damage | High, data would be exposed |
| Reproducibility | Low. Database is not directly exposed, authentication is in place |
| Exploitability | Low. Attacker should compromise at least another system first |
| Affected users | All |

Discoverability Low

Scientists stealing measurements

| Type | Level |
|-----------------|--|
| Damage | Medium. Peer reviews may be invalid |
| Reproducibility | Low. It requires another violationLow. It requires another violation |
| Exploitability | Low. It requires another violation |
| Affected users | Scientists involved in the experiments, external stakeholders |
| Discoverability | Low |

Denial of service

DDoS on APIs

| Type | Level |
|-----------------|---|
| Damage | High, system may become inoperative |
| Reproducibility | Low. The system should be exposed only in the internal network |
| Exploitability | Low. It would be easy to block the attack in the internal network |
| Affected users | All |
| Discoverability | Low. It would be difficult to plan an effective attack. |

DDoS on *Audit and Monitoring*

| Type | Level |
|-----------------|---|
| Damage | Low to High. It may cover a more vast attack |
| Reproducibility | Low. The system should be exposed only in the internal network |
| Exploitability | Low. It would be easy to block the attack in the internal network |
| Affected users | Auditors |
| Discoverability | Low. It would be difficult to plan an effective attack. |

Elevation of privilege

Scientists becoming administrators

| Type | Level |
|------|-------|
|------|-------|

| | |
|-----------------|--|
| Damage | High, the attacker could disrupt the system |
| Reproducibility | Low. It would require database access since no system function manipulates roles |
| Exploitability | Low. Attacker should compromise at least another system first |
| Affected users | All |
| Discoverability | Low |

System Requirements

Storage space

User and experiment data will require less than 1Kb per record, and their number is expected to be in the range of thousands. Therefore, it is safe to assume that a few megabytes will be sufficient to store them.

Each measurement is expected to require at least 22 bytes. With 1 million measures per experiment, each experiment will require about 21Mb of space.

| field | type | size |
|---------------|---------------------------------|---------|
| Measure type | integer | 2 bytes |
| Timestamp | Timestamp with nano precision | 8 bytes |
| Experiment id | integer | 4 bytes |
| Measure | Double precision floating point | 8 bytes |

CPU and memory

CPU and memory requirements will be determined with load testing after the initial deployment. Minimum resources will be set to values able to sustain the expected average daily traffic. Maximum resources will be set to values able to sustain 200% of the maximum expected traffic. Autoscale will be configured to follow the demand and contain costs.

GDPR Consideration

The application design requires only a minimal amount of personal information. All users will be able to retrieve, update and delete their own information, in compliance with GDPR. Complete deletion will preserve Staff Identification for traceability (TBD ref)

An administrator will be able to assist users with their GDPR request.

| Document | Field | Description |
|---------------|-----------------------------|--|
| User's record | Staff identification | Unique id number from the HR system |
| User's record | Name | Given name(s) |
| User's record | Surname | Family name |
| User's record | Email Address | Professional email address |
| Experiment | None | |
| Measurement | None | |
| Audit | User's staff identification | Only the user's staff identification will be stored in the audit |

References

- Aimar, A., Corman, A. A., Andrade, P., Fernandez, J. D., Bear, B. G., Karavakis, E., ... & Magnoni, L. (2019). MONIT: monitoring the CERN data centres and the WLCG infrastructure. In EPJ Web of Conferences (Vol. 214, p. 08031). EDP Sciences. Available from https://www.epj-conferences.org/articles/epjconf/pdf/2019/19/epjconf_chep2018_08031.pdf [Accessed 30 March 2022]
- Souza, R. D. A. (2020). Performance analysis between Apache Kafka and RabbitMQ. Available from <http://dspace.sti.ufcg.edu.br:8080/jspui/bitstream/riufcg/20339/1/RONAN%20DE%20ARAU%CC%81JO%20SOUZA%20-%20TCC%20CIE%CC%82NCIA%20DA%20COMPUTAC%CC%A7A%CC%83O%202020.pdf> [Accessed 30 March 2022]
- Khasawneh, T. N., AL-Sahlee, M. H., & Safia, A. A. (2020, April). Sql, newsql, and nosql databases: A comparative survey. In *2020 11th International Conference on Information and Communication Systems (ICICS)* (pp. 013-021). IEEE. Available from https://www.researchgate.net/profile/Mahmoud-Alsahlee/publication/340978543_SQL_NewSQL_and_NOSQL_Databases_A_Comparative_Survey/links/5ec46445a6fdcc90d685d608/SQL-NewSQL-and-NOSQL-Databases-A-Comparative-Survey.pdf [Accessed 30 March 2022]
- Kubernetes (N.D.) Horizontal Pod Autoscaling. Available from <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/> [Accessed 30 March 2022]

- OWASP (N.D) REST Security Cheatsheet. Available from https://cheatsheetseries.owasp.org/cheatsheets/REST_Security_Cheat_Sheet.html [Accessed 30 March 2022]

Bibliography (TBD remove)

- W3Techs (2022) Usage statistics of Nginx. Available from <https://w3techs.com/technologies/details/ws-nginx> [Accessed 30 March 2022]
- Rabbit: <https://www.rabbitmq.com/production-checklist.html>