

MyMONIT

Collecting measurements to monitor CERN's experiments – README

Team 2: Ahmed, Padukka, Rossotto, Thompson, Wimalendran

Table of Contents

Application composition.....	3
Overview.....	3
Components.....	3
Class diagram of APP.....	6
Overview.....	6
RestListener.....	6
RabbitListener.....	7
Logging.....	8
Storage.....	9
Dependency Injection.....	10
Orchestration.....	13
Technical specifications.....	15
Requirements.....	15
Containers overview.....	15
Summary of container images.....	15
Production configuration.....	16
Scalability of App.....	16
Scalability of MySQL, Nginx, and RabbitMQ.....	16
Scalability of Logstash, Elasticsearch, and Kibana.....	17
Ports.....	17

Database scripts and default data.....	17
Automatic build.....	18
Authentication.....	19
User creation.....	19
Authentication process.....	20
Password policy.....	21
How to get a token.....	21
Testing.....	23
Unit testing.....	23
Automatic scanners.....	23
Bandit.....	24
Coverage.....	24
Pylint.....	25
Manual testing.....	25
APIs sequence.....	26
Calling the APIs.....	27
API documentation.....	28
Differences in the specifications.....	29
List of attachments.....	32
Repository.....	32
References.....	33

Application composition

Overview

As mentioned in the analysis (Rossotto et al, 2022), CERN uses a variety of independently developed systems to monitor the infrastructure used in the experiments (Aimar et al., 2019). MyMONIT unifies the monitoring and integrates the streams of different applications into a single repository.

MyMonit collects the input from the experiments (the measures) via message broker and exposes REST APIs to consume them. MyMonit exposes a separate interface for auditing the solution itself. The whole solution is composed of different containers orchestrated by Docker Compose.

Components

The application is composed of seven containers running in Docker.

- App is the core component. It hosts the business logic and it is written in Python. It exposes REST APIs and consumes messages via AMQP. Its stack is composed of
 - Flask (Flask, N.D.): a framework used to create REST APIs.
 - Pika (Pika, N.D.): a framework to consume messages from the AMQP interface.
 - MySQL Connector (MySQL Connector, N.D.): database connector.
 - Python Logstash (Elastic, N.D.): connector to the Logstash server to store logs for auditing purposes.

- Firebase Admin (Firebase, N.D.): connector to Firebase for the authentication of JSON Web Tokens.
- Dependency Injector (Dependency Injector, N.D.): a framework to implement dependency injection in Python.
- Colorama (Hartley J. (2021): a framework to output ANSI colored text in console.
- MySQL (MySQL, N.D.) is the core database. It contains the users with access to App, and the data relative to the experiments.
- RabbitMQ (RabbitMQ, N.D.) is the AMQP message broker responsible for the collection of the measures and the delivery to App.
- Nginx (Nginx, N.D.) is the HTTP/S proxy that exposes App's REST APIs to the external world.
- Logstash (Logstash, N.D.) is responsible for log collection, formatting, and correct storage in Elasticsearch.
- Elasticseach (Elasticsearch, N.D.) is the database responsible for the collection of all the logs from all other containers.
- Kibana (Kibana, N.D.) is the frontend application that visualizes the content of Elasticsearch.

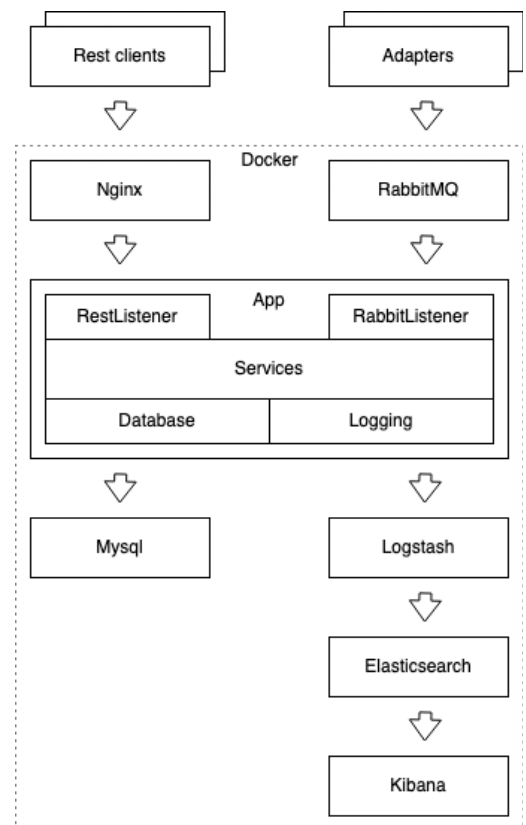


Figure 1

Additionally, the application supports "adapters": small applications designed to translate external input into a correctly formatted message that can be enqueued in RabbitMQ.

Figure 1 illustrates the overview of the components and their interaction.

Class diagram of APP

Design principles

The modular design aims to follow SOLID principles (Martin, R, 2000) with particular reference to single responsibility, open-closed principle, and dependency inversion.

Model

The model is composed of Measure, Experiment, User, and Exceptions (Figure 2).

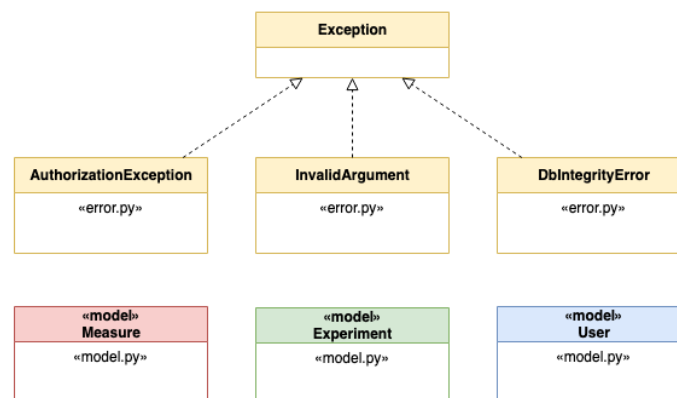


Figure 2

Three specialised services are responsible for the three main elements of the model (Figure 3).

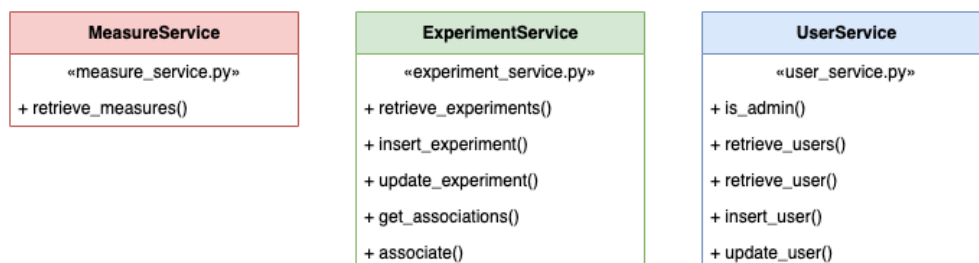


Figure 3

RestListener

RestListener is responsible for the REST interface of the application. RestListener uses RestConfiguration to model the configuration, and delegates the handling of the resources to specialised resource components (MeasureResource, ExperimentResource, ExperimentAssociationResource, UserResource, and ExceptionResource) (Figure 4).

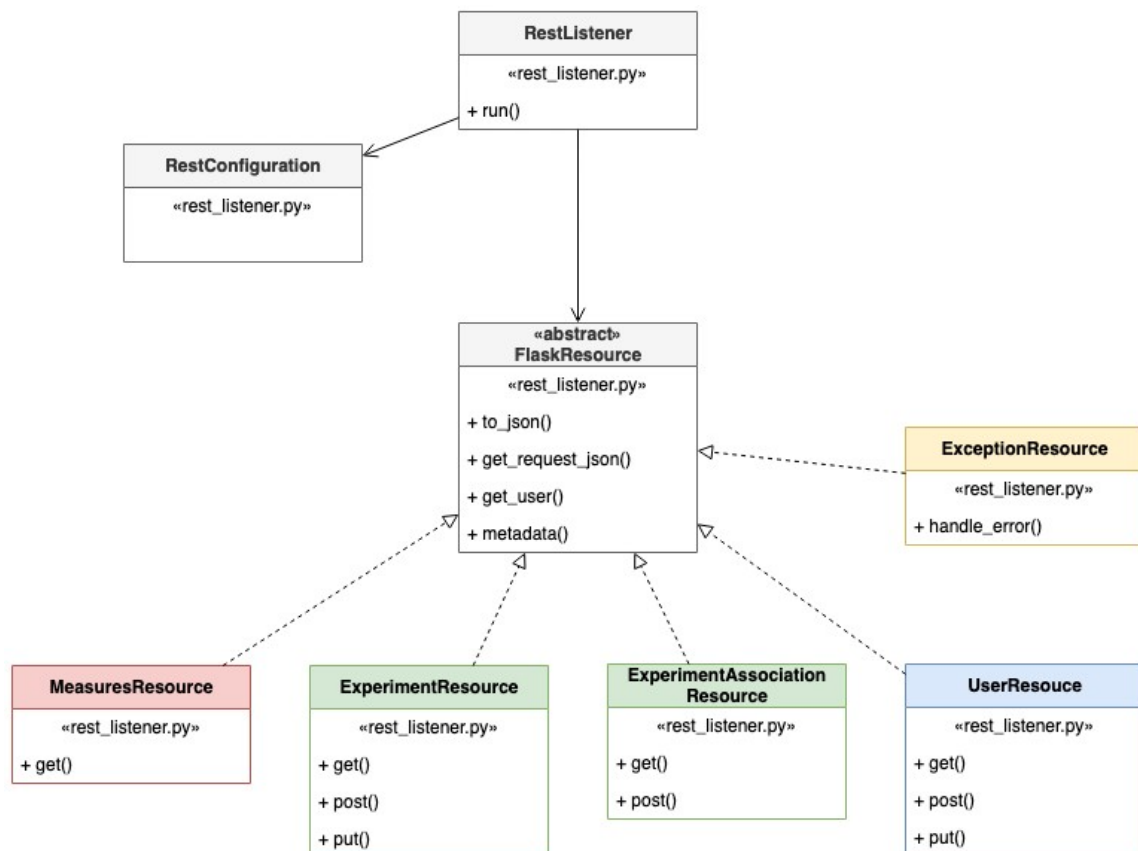


Figure 4

RabbitListener

RabbitListener is the service responsible for the connection to RabbitMQ and it uses a RabbitConfiguration to model the configuration. RabbitConnector is a component of the listener used to extract the IO functions and simplify unit testing.

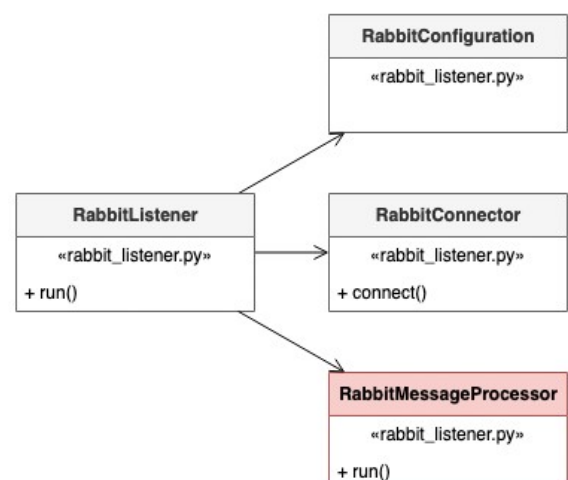


Figure 5

The extraction makes it possible to mock the component and avoid a real I/O in the test. RabbitMessageProcessor contains the logic to store the measures in the database (Figure 5).

Logging

Logging is a cross-cutting service used in different components.

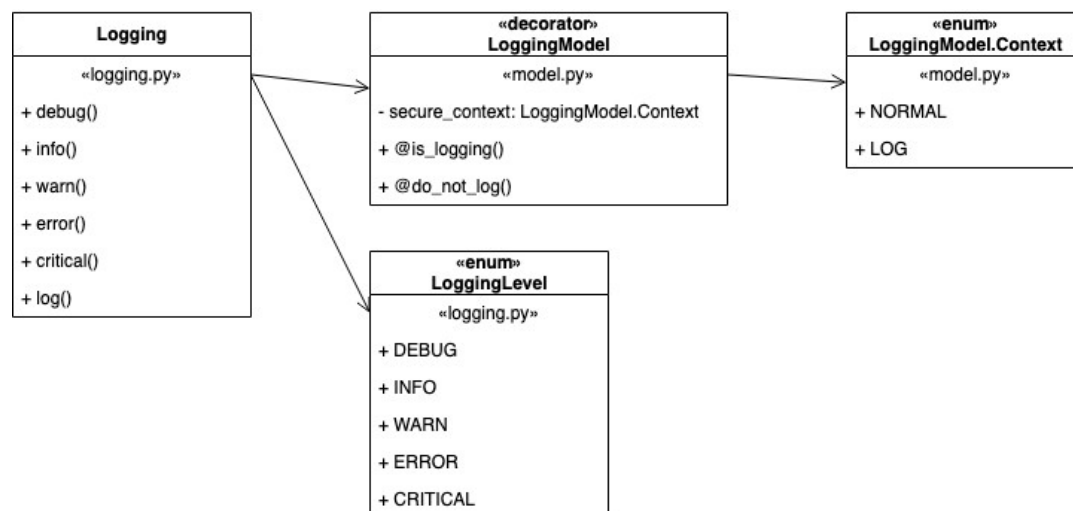


Figure 6

Logging supports faceted data (Schmitz, 2016) via two decorators `@is_logging` and `@do_not_log` to respectively mark functions that perform logging, and the parts of the model that should not appear in the logs with their real value to prevent data leaks. `@is_logging` wraps the decorated call setting the context's value (Figure 7).

```

@staticmethod
def is_logging(func):
    def inner(*args, **kwargs):
        LoggingModel._set_context_value(LoggingModel.Context.LOG)
        func(*args, **kwargs)
        LoggingModel._set_context_value(LoggingModel.Context.NORMAL)
    return inner
  
```

Figure 7

@do_no_log checks the context's value and redirects the call to the original method or returns "***" depending on the context's value. This simple mechanism can hide secrets in the logs without implementing a case-by-case logic in the logging component (Figure 8).

```
@staticmethod
def do_not_log(func):
    def inner(*args, **kwargs):
        return func(*args, **kwargs) if LoggingModel._get_context_value() == \
            LoggingModel.Context.NORMAL else '***'
    return inner
```

Figure 8

The component supports multithreading with a dictionary of contexts indexed by thread id (Figure 9).

```
@staticmethod
def _get_context():
    return LoggingModel.Context.secure_context

@staticmethod
def _get_context_value():
    thread_id = threading.get_native_id()
    context = LoggingModel._get_context()
    if thread_id not in context:
        context[thread_id] = LoggingModel.Context.NORMAL
    return context[thread_id]

@staticmethod
def _set_context_value(value:Context):
    thread_id = threading.get_native_id()
    context = LoggingModel._get_context()
    context[thread_id] = value
```

Figure 9

Storage

Storage is responsible for the IO with the database (Figure 10).

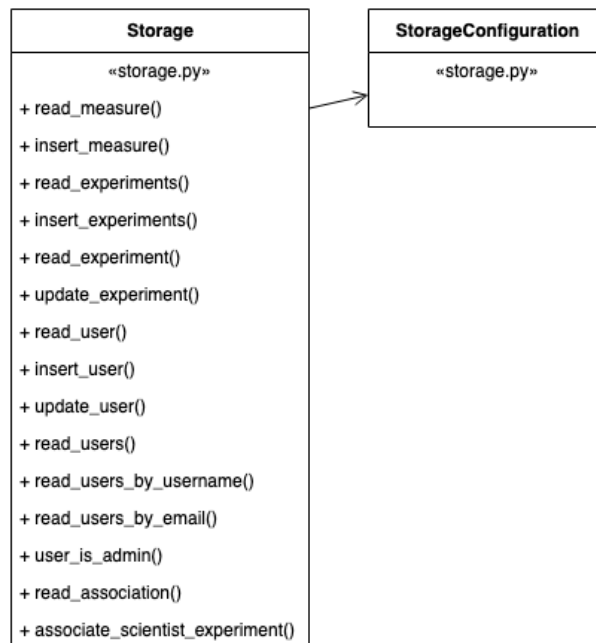


Figure 10

Dependency Injection

Container is responsible for the configuration and dependency injection. Thanks to Container, components do not instantiate one another simplifying unit testing (Figure 11).

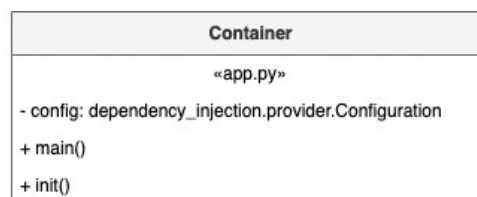


Figure 11

The configuration of Container is in `app.py`. The parameters have a default value that can be overridden by the environment's parameters (Figure 12).

```
def init():
    '''initialise the application'''
    container = Container()

    # db
    container.config.db_user.from_env('DB_USER', default = 'root', as_ = str)
    container.config.db_password.from_env('DB_PASSWORD', default = 'password', as_ = str)
    container.config.db_host.from_env('DB_HOST', default = 'localhost', as_ = str)
    container.config.db_database.from_env('DB_DATABASE', 'my_monit', as_ = str)
    # rabbit
    container.config.rabbit_url.from_env('RABBIT_URL', default = 'localhost', as_ = str)
    container.config.rabbit_user.from_env('RABBIT_USER', default = 'guest', as_ = str)
    container.config.rabbit_password.from_env('RABBIT_PASSWORD', default = 'guest', as_ = str)
    container.config.rabbit_exchange.from_env('RABBIT_EXCHANGE', default = 'mymonit', as_ = str)
    container.config.rabbit_routing.from_env('RABBIT_ROUTING', default = 'measures', as_ = str)
    container.config.rabbit_queue.from_env('RABBIT_QUEUE', default = 'measures', as_ = str)
    # rest
    container.config.rest_host.from_env('REST_HOST', default = '0.0.0.0', as_ = str)
    # logstash
    container.config.logstash_host.from_env('LOGSTASH_HOST', default = 'localhost', as_ = str)
    container.config.logstash_port.from_env('LOGSTASH_PORT', default = 5959, as_ = int)
    # firebase
    if 'GOOGLE_APPLICATION_CREDENTIALS' not in os.environ:
        # os.getcwd() is the whole project's root
        os.environ['GOOGLE_APPLICATION_CREDENTIALS'] = './containers/app/private_key.json'
```

Figure 12

The dependency injection mechanism in Container uses the `containers.config` to create the objects and inject them one into another.

For example, logging is a basic object without dependencies and it is injected into storage (Figure 13).

Dependency injection is essential to perform accurate unit testing because it makes object-mocking easy.

```
logging = providers.Singleton(
    Logging,
    host = config.logstash_host,
    port = config.logstash_port
)

storage_configuration = providers.Singleton(
    StorageConfiguration,
    db_user = config.db_user,
    db_password = config.db_password,
    db_host = config.db_host,
    db_database = config.db_database
)

rabbit_configuration = providers.Singleton(
    RabbitConfiguration,
    rabbit_url = config.rabbit_url,
    rabbit_user = config.rabbit_user,
    rabbit_password = config.rabbit_password,
    rabbit_exchange = config.rabbit_exchange,
    rabbit_routing = config.rabbit_routing,
    rabbit_queue = config.rabbit_queue
)

rest_configuration = providers.Singleton(
    RestConfiguration,
    rest_host = config.rest_host
)

storage_connector = providers.Singleton(
    StorageConnector,
    storage_configuration = storage_configuration,
    logging = logging
)

storage = providers.Singleton(
    Storage,
    storage_connector = storage_connector,
    logging = logging
)
```

Figure 13

Orchestration

The orchestration uses Docker Compose (Docker Compose, N.D.) via the descriptor `docker-compose.yml` (Figure 14).

```
app:
  image: safe_repository:latest
  container_name: app
  environment:
    - DB_USER=root
    - DB_PASSWORD=password
    - DB_HOST=mysql
    - DB_DATABASE=my_monit
    - RABBIT_URL=rabbit
    - RABBIT_USER=guest
    - RABBIT_PASSWORD=guest
    - RABBIT_EXCHANGE=mymonit
    - RABBIT_ROUTING=measures
    - RABBIT_QUEUE=measures
    - REST_HOST=0.0.0.0
    - LOGSTASH_HOST=logstash
    - LOGSTASH_PORT=5959
    - GOOGLE_APPLICATION_CREDENTIALS=/firebase_config.json
  depends_on:
    - elasticsearch
    - rabbit
    - mysql
  volumes:
    - ./containers/app/private_key.json:/firebase_config.json
```

Figure 14

The environment list specifies the variables for the container, including the secrets. The included example does not use strong secrets and it is supposed to be only an example. It is recommended to change the secrets in a production environment. It is also recommended to use this configuration as a model to deploy into Kubernetes that offers a stronger solution for deployment, including the possibility to separate the secrets from the code (Kubernetes, N.D.).

Some ports, such as the RabbitMQ console on 15672 (Figure 15) are exposed to easily demonstrate the usage of the components. Only the ports mentioned in the section “production configuration” are required for correct functioning.

```
rabbit:
  image: "rabbitmq:3-management"
  container_name: rabbit
  ports:
    - "5672:5672"
    - "15672:15672"
```

Figure 15

Technical specifications

Requirements

It is required to have the Docker daemon running. For the development, the chosen platform was Docker Desktop. It is recommended to assign at least 4gb of ram to Docker: with less memory, some containers may fail. It is always possible to restart failing containers should that happen.

All scripts require a unix-like system (eg. Mac).

Python 3.9 is required to build App alone. All dependencies are specified in setup.py and can be installed with “pip3 install -e .”

Containers overview

The configuration of all containers is visible in “docker-compose.yml”. The file is meant to be an example and it is not suitable for the production environment because it contains unencrypted secrets and exposes internal ports.

Summary of container images

Container's name	Base image	Customisation
App	Python:3.9	Injection of the application
Nginx	Nginx:latest	Inject of configuration and SSL certificates
Rabbit	Rabbit:3-management	None

Container's name	Base image	Customisation
Mysql	Mysql:5.7	None
Elasticsearch	Elasticsearch:6.5.4	None
Kibana	Kibana:6.5.4	None
Logstash	Logstash:6.5.4	None

Production configuration

It is recommended to deploy the application in Kubernetes. Kubernetes makes it easy to share secrets safely without embedding them in the code and offers advanced functions to ensure scalability.

Scalability of App

The core Python application is stateless and it is possible to start multiple instances in parallel to improve performance.

If the goal is to increase the performance of the REST APIs, a change in the Nginx configuration is required to implement load-balancing between the nodes.

If the goal is to increase performances of the consumption of AMQP messages through the RabbitMQ interface, it is sufficient to increase the number of nodes to multiply the consumers. All the calls to the REST APIs will be handled by one single node, but this does not represent an issue.

Scalability of MySQL, Nginx, and RabbitMQ

These components natively support clustering through configuration (MySQL – Clustering, N.D.; Nginx – Clustering, N.D; RabbitMQ – Clustering, N.D.).

Scalability of Logstash, Elasticsearch, and Kibana

The ELK stack supports clustering through configuration (Elastic – Clustering. N.D.).

It is reasonable to expect that the components under pressure may be Logstash and Elasticsearch because they have the responsibility to store logs from all components. It is an assumption that Kibana's users will be limited, therefore kibana is not expected to suffer from scalability problems.

Ports

The only ports that need to be published in a production environment are:

- 443: Nginx
- 5672: RabbitMQ
- 5601: Kibana

All other ports in the docker-compose descriptor are exposed only for testing and should remain closed to external connections in a production environment.

Database scripts and default data

The MySQL container is preconfigured with a schema. All the definitions can be found in `containers/mysql/dbschema`.

For evaluation purposes, the folder with the schemas also contains examples of data.

The examples include:

- Three users
 - One administrator
 - Two scientists

- Two experiments
- Sample measures

All the examples are injected with the script “999_default_data.sql” which should be removed before the installation in a production environment.

Automatic build

The script “run.sh” automatically builds and starts all the containers.

The script “run-infra.sh” starts only a subset of the containers (Mysql and RabbitMQ) and it is convenient for the development of the core application in Python.

The Python application can be built with “pip3 install -e .” launched in the App’s container base folder. The command downloads automatically the dependencies specified in “setup.py”.

Authentication

The authentication process is externalized. The application contains GDPR-sensitive data (like emails), but no credentials in any form.

The chosen platform for the external authentication is Firebase (Firebase, N.D.) hosted in the cloud by Google.

Users must register and verify their identity only with Firebase to generate a signed JSON Web Token that contains their unique id and can be verified by App during a REST call.

User creation

A user must be created in the Firebase console (Firebase – Console, N.D.). The console can be accessed using the credentials in the attachment.

The console shows all the users and their User UID which is the information that links a Firebase's User with App's User (Figure 16).

Authentication

[Users](#)[Sign-in method](#)[Templates](#)[Usage](#)

Search by email address, phone number, or user UID

Add user

Identifier	Providers	Created ↓	Signed In	User UID
kbouman@home.cern		May 5, 2022	May 6, 2022	hirbBCceXIWUgdpNHUpaNtworjm2
phigg@home.cern		May 5, 2022	May 14, 2022	uJayGH994iVaZTY5auSNYZoUpw...
aalcorn@home.cern		May 5, 2022	May 13, 2022	5MS0kr1nE5Qm7ZNzAyHNCVfap...
padukka.sathira@gmail.com		May 3, 2022	May 3, 2022	MIUWT1Prh6boAsUB0eihfFVERW...
sathira.padukka@gmail.com		May 2, 2022	May 12, 2022	zlp88fmXEIZ8ygj7yOgsZFo3uPt2
alberto.rossotto@gmail.com		Apr 30, 2022	May 5, 2022	j2GDTte91eYUcuHmgw0D9ylOvX...
testingperson2496@gmail....		Apr 30, 2022	May 5, 2022	sexV7ZPFOxYOWLxvpdyDSADvoT...

Rows per page: 50

1 – 7 of 7

Figure 16

If a new user needs to be created, the user can navigate to the console page, enter the credentials, and register the user. This will create a new user in Firebase.

The corresponding user must be then created in APP using the User API (Figure 17).

```
curl --location --request POST 'http://mymonit/users/' \
--header 'Content-Type: application/json' \
--header 'Authorization: Bearer ...' \
--data-raw '{
  "user_id": "S123",
  "name": "Jim Apple",
  "username": "firebaseUid",
  "email": "jim@foobar.com",
  "role": "SCIENTIST"
}'
```

Figure 17

The field “username” must correspond to the User’s uid in Firebase.

Authentication process

The authentication generates a signed Json Web Token with the user’s information on the Firebase side. The token is included as Bearer token to the App’s backend in each

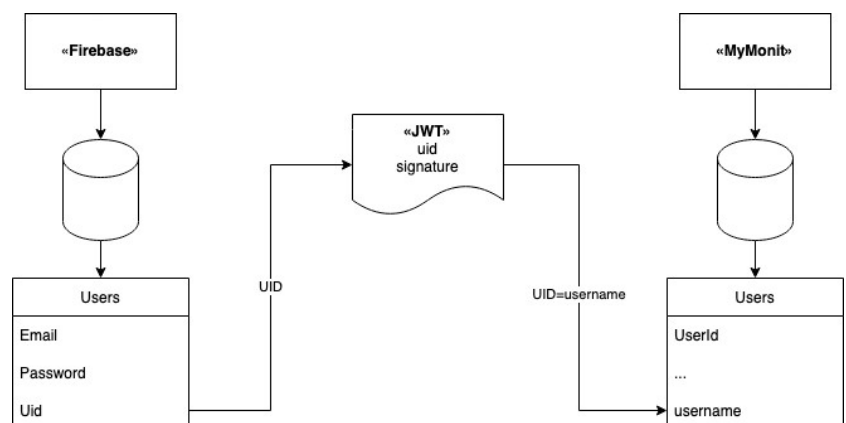


Figure 18

REST API call. App’s backend verifies the signature and maps the Firebase’s Uid with App’s username. App can then use the authenticated user for the authorization of each call (Figure 18).

Password policy

This aspect is not implemented in MyMonit. It is the responsibility of the administrators of the authentication service (Firebase, in this case) to enforce a password policy.

How to get a token

Open with a browser {url}/static/index.html

A login and registration form with two input fields and two buttons. The first input field is labeled "Enter your email" and the second is labeled "Enter your password". Below the second input field are two buttons: "Sign In" and "Register".

Enter your email
Enter your password
Sign In
Register

Figure 19

Fill the credentials in the form (Figure 19) and click “sign in”. If the authentication is successful, the next page will show the JWT that must be included in the REST calls (Figure 20).



aalcorn@home.cern

eyJhbGciOiJSUzI1NiIsImtpZCI6ImJlYmYxMDBlYWRRkYTMzMmVAT34lBaSkCVUFHR0ElJXyJW8JH53n8Jq_zvg80_9z42e9CxxqICYxqXxX3r4DQnnskRremJgrevOmg0aouTF7GrSfXLbdjn78qFDbZ2I0roN

Figure 20

Using Postman, the tokens can be set as variables (Figure 21).

Testing

Unit testing

Unit tests cover a vast portion of the application to help detect regressions during the development.

Unit test use of Mock and MagicMock (Python Mock, N.D.) to replace the dependencies in the class under test. The mocks can return values set programmatically and make it easy to test how the class under test called them.

```
def test_insert_experiment_ok(self):  
    '''verify insert'''  
    self.user_service.is_admin = MagicMock(return_value = True)  
    self.experiment_service.insert_experiment({'name': 'abc'}, '1234')  
    self.storage.insert_experiment.assert_called_with(Experiment(experiment_id = None,  
                                                                name = 'abc'))
```

Figure 22

In the test in Figure 22, `user_service` and `storage` are mocked. `MagicMock` ensures that the call to `is_admin` returns `True` in order to pass an authorization check, and the test verifies that the class under test, `experiment_service`, calls `insert_experiment` on `storage` with specific parameters.

Automatic scanners

The project uses `Pylint` (Pylint, N.D.) and `Bandit` (Python Code Quality Authority, 2022) to spot unused variables, formatting issues, and potential vulnerabilities. `Coverage.py` (Coverage.py, N.D.) measures the code coverage. It is possible to

launch the three utilities with the run-checks.sh script in containers/app with results in containers/app/reports. The current test results are available as attachments.

Bandit

Bandit reports a possible vulnerability:

```
>> Issue: [B104:hardcoded_bind_all_interfaces] Possible binding to all  
interfaces.
```

Severity: Medium Confidence: Medium

This refers to the default configuration of the APP container that binds Flask on 0.0.0.0. It is not a vulnerability: the restriction on the connectivity is delegated to Docker.

Coverage

The test coverage is above 80%. The report shows precisely which lines are not covered, and they are mostly lines that would not be useful to test (eg. getter) or it would require heavy usage of

Coverage report: 81%				
Module	statements	missing	excluded	coverage
my_monit/__init__.py	0	0	0	100%
my_monit/app.py	57	0	0	100%
my_monit/errors.py	9	3	0	67%
my_monit/experiment_service.py	45	6	0	87%
my_monit/logging.py	48	2	0	96%
my_monit/measure_service.py	12	0	0	100%
my_monit/model.py	115	7	0	94%
my_monit/rabbit_listener.py	71	12	0	83%
my_monit/rest_listener.py	157	49	0	69%
my_monit/storage.py	107	42	0	61%
my_monit/user_service.py	50	6	0	88%
Total	671	127	0	81%

Figure 23

mocking rendering the test insignificant (eg. I/O) (Figure 23).

Pylint

The report mentions one possible defect:

```
***** Module my_monit.logging
```

```
my_monit/logging.py:78:15: W0703: Catching too general exception Exception  
(broad-exception)
```

Your code has been rated at 9.98/10

It refers to a line in a method that performs logging. Although it is not a good practice, catching Exception aims to capture any unforeseen issue during logging and print it in stdout as a fallback.

Manual testing

Manual testing was used to quantify the differences between the initial specifications and the implementation, and ensured the overall quality of the application.

The following are useful URLs to inspect the application when deployed with the provided test configuration.

Service	URL
Base APIs URL	https://localhost
RabbitMQ Console	http://localhost:15672/
Kibana	http://localhost:5601/login
Mysql Console	localhost:3306

APIs sequence

The following diagram illustrates a valid sequence of calls.




Figure 24

An administrator must create a User of type SCIENTIST. Subsequently, the ADMIN can retrieve and modify the User. Similarly, the Admin can create an Experiment.

With at least one User and one Experiment, it is possible to associate them to give access to the Experiment.

When associated, the User can retrieve the Measurements.

An ADMIN cannot read the Measurements, and a SCIENTIST has limited or no access to most of the APIs.

Calling the APIs

A collection of all the API calls is available for Postman (Postman, N.D.) in the attachments (Figure 25).

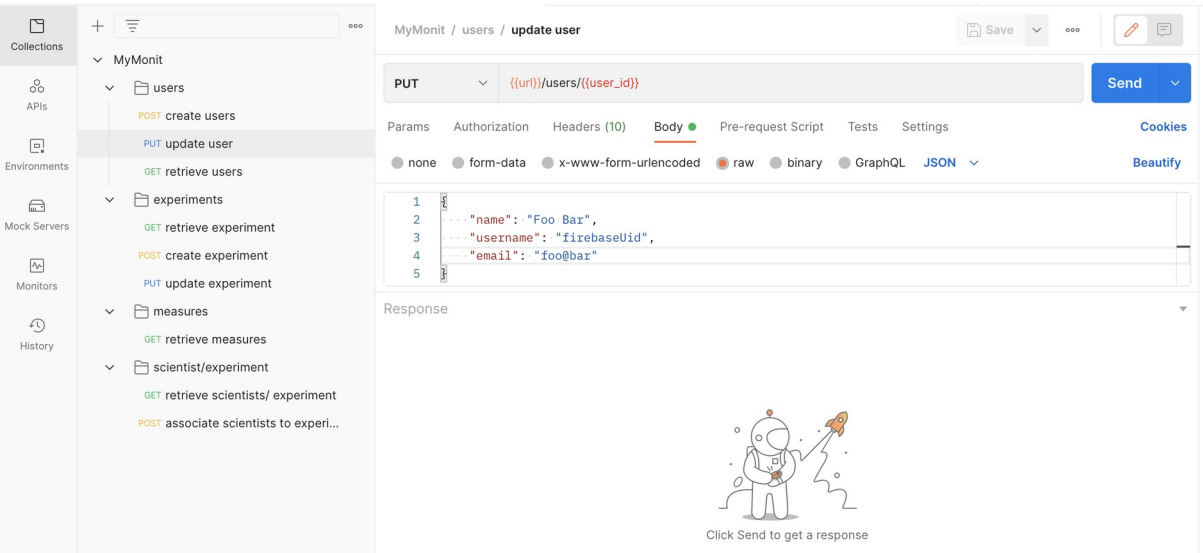


Figure 25

The collection uses variables that can be specified in the Environment tab (Figure 26).

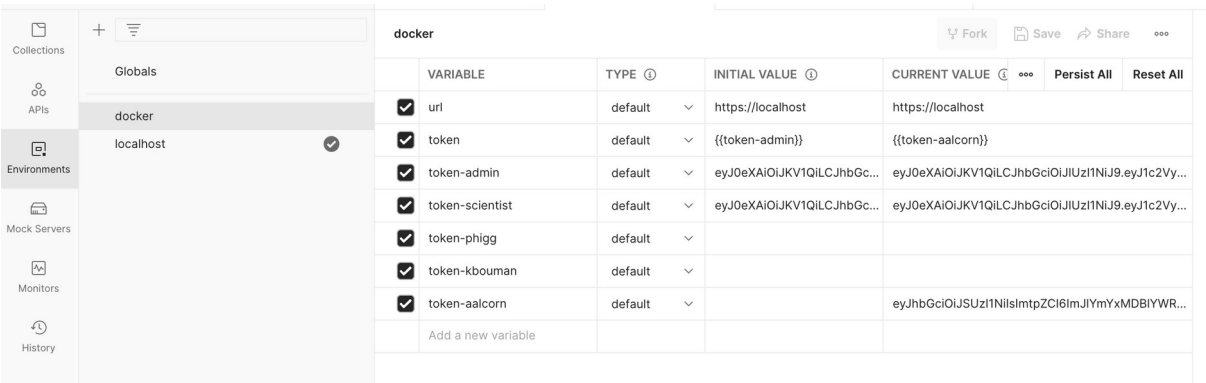


Figure 26

API documentation

The APIs are documented in the attached api.html using Swagger (Swagger, N.D.) (Figure 27).

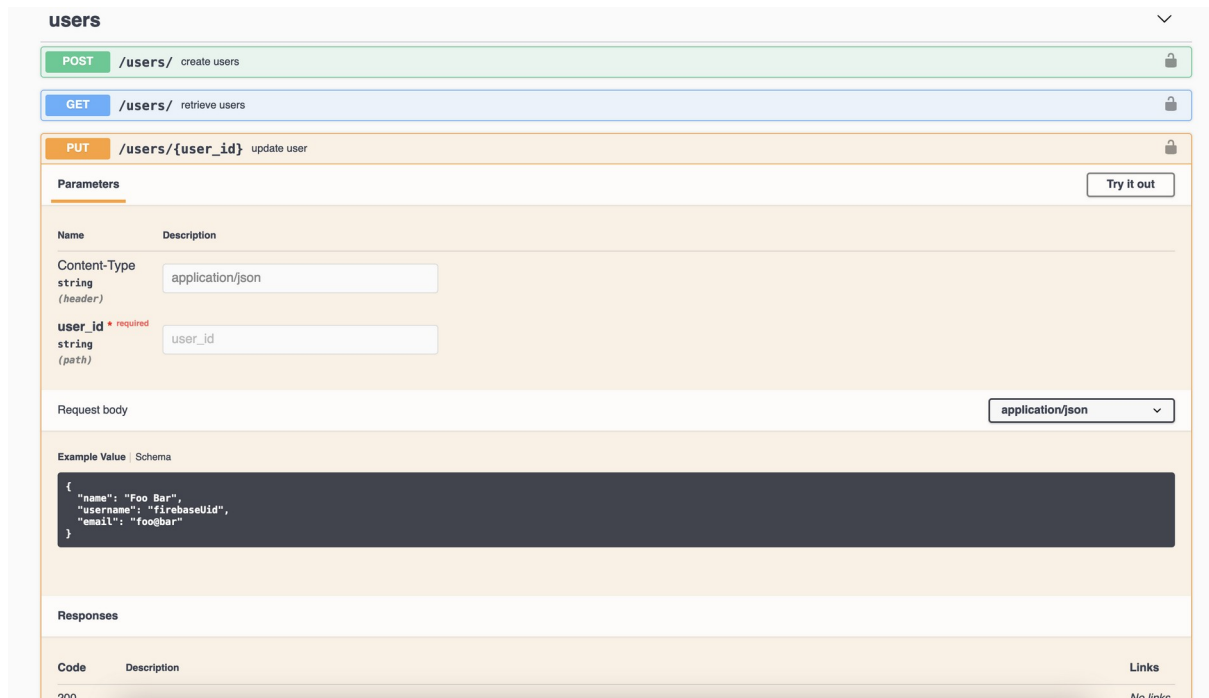


Figure 27

Differences in the specifications

The following is an adaptation of the original use case diagram. It illustrates which use-cases were fully implemented (in green), partially implemented (in yellow), and omitted from the final deliverable (in red).




Figure 28

- Adapters: only one adapter was developed as a proof of concept to demonstrate the functionality.
- User management: it is not currently possible to explicitly delete users. The deletion of the user was delayed because it cannot be a simple deletion. Part of the user's record must be preserved for auditing, while all information non strictly required for auditing must be removed in compliance with GDPR. It is however possible to update the user to achieve the same practical result, although the solution is not ideal. It is also observed that no information that allows the direct identification of a user is being logged, therefore auditing is not affected by the issue.
- RabbitMQ encryption: it was not implemented. It is, therefore, possible to intercept or manipulate measurements unless other technical solutions (eg. a VPN) are in place.
- Autoscalability: this requirement was ignored because Docker Compose does not support the functionality. It is however possible to run the application in Kubernetes and configure autoscale linked to CPU usage (Kubernetes – autoscale, N.D.).
- Log event: this requirement was implemented only partially.
 - App's logs are transmitted to Logstash directly and are available in Kibana. This requirement is fully implemented.
 - Nginx's logs are transmitted to Logstash with Filebeat (Filebeat, N.A.) and are available in Kibana. Filebeat works as an adapter between Nginx's

logs and Logstash. This requirement is fully implemented. This requirement is fully implemented.

- MySQL's logs are available only locally through Docker. The implementation of the connection to the ELK stack suffered from a configuration issue that made it unstable. It was impossible to connect MySQL directly to Logstash. There was an attempt to use Filebeat but the service could not start in the correct order inside the container requiring the manual intervention of an operator to establish a connection. Problems persisted also with manual intervention because the volume of logs caused stability issues in Elasticsearch. This requirement is not implemented.
- RabbitMQ logs: they are available only locally through Docker. The implementation suffered from the same issues as the MySQL's logs. This requirement is not implemented.

List of attachments

- Firebase: credentials for the Firebase console and for the default users. This attachment contains credentials and should be removed when the documentation is shared with third parties.
- Postman: postman collection with examples of the APIs usage.
- Reports: reports generated with 'run-check.sh'
 - bandit.txt: a report generated with bandit
 - cov: a report generated with coverage
 - pylint.txt: a report generated with pylint
- Swagger: description of the APIs

Repository

The project's repository is available at <https://github.com/ros101/ssdcs-assignment>

References

- Aimar, A., Corman, A. A., Andrade, P., Fernandez, J. D., Bear, B. G., Karavakis, E., ... & Magnoni, L. (2019). MONIT: monitoring the CERN data centres and the WLCG infrastructure. In EPJ Web of Conferences (Vol. 214, p. 08031). EDP Sciences. Available from https://www.epj-conferences.org/articles/epjconf/pdf/2019/19/epjconf_chep2018_08031.pdf [Accessed 14/May/2022]
- Coverage.py (N.D.) Coverage.py. Available from <https://coverage.readthedocs.io/en/6.3.3/> [Accessed 14/May/2022]
- Dependency Injector (N.D.) Dependency Injector — Dependency injection framework for Python. Available from <https://python-dependency-injector.ets-labs.org/> [Accessed 14/May/2022]
- Docker Compose (N.D.) Overview of Docker Compose. Available from <https://docs.docker.com/compose/> [Accessed 14/May/2022]
- Elastic (N.D.) Centralize, transform & stash your data. Available from <https://www.elastic.co/logstash/> [Accessed 14/May/2022]
- Elasticsearch (N.D.) What is Elasticsearch? Available from <https://www.elastic.co/what-is/elasticsearch> [Accessed 14/May/2022]
- Elastic (N.D.) Set up a cluster for high availability. Available from <https://www.elastic.co/guide/en/elasticsearch/reference/current/high-availability.html> [Accessed 14/May/2022]

- Filebeat (N.D.) Lightweight shipper for logs. Available from <https://www.elastic.co/beats/filebeat> [Accessed 14/May/2022]
- Firebase (N.D.) Firebase. Available from <https://firebase.google.com/> [Accessed on 14/May/2022]
- Firebase (N.D.) Add the Firebase Admin SDK to your server. Available from <https://firebase.google.com/docs/admin/setup> [Accessed 14/May/2022]
- Firebase (N.D.) Firebase Console. Available from <https://console.firebase.google.com/project/coding-output/authentication/users> [Accessed 14/May/2022]
- Flask (N.A.) Flask, web development one drop at a time. Available from <https://flask.palletsprojects.com/en/2.1.x/> [Accessed 14/May/2022]
- Hartley J. (2021) Colorama. <https://github.com/tartley/colorama> [Accessed 14/May/2022]
- Kibana (N.D.) Your window into the Elastic Stack. Available from <https://www.elastic.co/kibana/> [Accessed 14/May/2022]
- Kubernetes (N.D.) Secrets. Available from <https://kubernetes.io/docs/concepts/configuration/secret/> [Accessed 14/May/2022]
- Kubernetes (N.D.) Horizontal Pod Autoscaling. Available from <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/> [Accessed 14/May/2022]

- Martin, R (2000) Available from https://web.archive.org/web/20150906155800/http://www.objectmentor.com/resources/articles/Principles_and_Patterns.pdf [Accessed 14/May/2022]
- MySQL Connector (N.D.) MySQL Connector/Python Developer Guide. Available from <https://dev.mysql.com/doc/connector-python/en/> [Accessed 14/May/2022]
- MySQL (N.D.) MySQL - The world's most popular open source database. Available from <https://www.mysql.com/> [Accessed 14/May/2022]
- MySQL (N.D.) MySQL Cluster: Getting Started. Available from <https://www.mysql.com/products/cluster/start.html> [Accessed 14/May/2022]
- Nginx (N.D.) Nginx. Available from <https://nginx.org/en/> [Accessed 14/May/2022]
- Nginx (N.D.) What Is Clustering? <https://www.nginx.com/resources/glossary/clustering/> [Accessed 14/May/2022]
- Pika (N.D.) Introduction to Pika. Available from <https://pika.readthedocs.io/en/stable/> [Accessed 14/May/2022]
- Pylint (N.D.) Pylint User Manual. Available from <https://pylint.pycqa.org/en/latest/> [Accessed 14/May/2022]
- Python Code Quality Authority (2022) Bandit. Available from <https://github.com/PyCQA/bandit> [Accessed 14/May/2022]

- Python Mock (N.D.) unittest.mock — mock object library. Available from <https://docs.python.org/3/library/unittest.mock.html> [Accessed 14/May/2022]
- Postman (N.D.) Postman. Available from <https://www.postman.com/> [Accessed 14/May/2022]
- RabbitMQ (N.D.) RabbitMQ. Available from <https://www.rabbitmq.com/> [Accessed 14/May/2022]
- RabbitMQ (N.D.) Clustering Guide. Available from <https://www.rabbitmq.com/clustering.html> [Accessed 14/May/2022]
- Rossotto, A, Ahmed, B, Padukka, S, Thompson, V, Wimalendran, P (2022) MyMONIT - Collecting measurements to monitor CERN's experiments.
- Schmitz, T., Rhodes, D., Austin, T.H., Knowles, K., Flanagan, C. (2016) Faceted Dynamic Information Flow via Control and Data Monads. In: Piessens F., Viganò L. (eds) Principles of Security and Trust. POST 2016. Lecture Notes in Computer Science, vol 9635. Springer, Berlin, Heidelberg.
- Swagger (N.D.) Swagger. Available from <https://swagger.io/> [Accessed 14/May/2022]